

Gradient Boosting Machine

A single decision tree is like a lone detective solving a mystery: swift to act but sometimes overconfident in early leads. Now imagine a team of detectives working sequentially – each reviewing the case, learning from previous mistakes, and sharpening the investigation. That's a **Gradient Boosting Machine**, a model that builds trees one after another, where each tree corrects the missteps of the last, gradually closing in on the truth.

In the previous post, we explored how Random Forests combine many independent trees grown in parallel to mitigate bias and variance. This time, we'll examine a different approach: **boosting**. Boosting is an ensemble learning technique that combines a series of weak learners into a strong learner, minimizing training errors by focusing increasingly on the hardest-to-predict data. Specifically, we'll explore the Gradient Boosting Regressor, which trains trees in an ordered sequence to produce highly accurate predictions.

What is Gradient Boosting?

Gradient Boosting builds an ensemble of weak learners, typically shallow decision trees, in a sequential manner. Each new model is trained to reduce the loss function of the ensemble by fitting the negative gradient (also known as pseudo-residuals) of the loss with respect to current predictions. This approach iteratively corrects errors made by previous models using gradient descent.

How does this differ from boosting in general? Traditional boosting methods like AdaBoost adjust the weights of misclassified samples to focus learning. Gradient boosting instead directly models the gradient of a loss function, fitting new learners to the residual errors, which are the differences between true values and predictions of previous models.

Gradient Boosting Machines (GBMs) belong to this family of algorithms, which build the ensemble stage-wise and continually refine predictions by focusing on residual errors. Both

regression and classification tasks benefit from Gradient Boosted Decision Trees; popular implementations include XGBoost, LightGBM, and CatBoost. Today, we'll focus on the Gradient Boosting Regressor applied to predicting Customer Lifetime Value (CLV). For context, CLV is the total revenue a customer generates over their engagement with a business.

Structure

A Gradient Boosting Machine assembles its predictive power from a sequence of decision trees, but unlike Random Forests, these trees are grown sequentially, not independently or in parallel. Each tree is designed specifically to address the weaknesses of the current ensemble, honing in on data points where errors remain high and enhancing prediction accuracy iteratively.

At the core of this process is gradient descent, an optimization technique that guides the model to minimize a chosen loss function. For regression problems, this is often the Mean Squared Error (MSE); for classification, cross-entropy loss or similar functions are used. The model begins with an initial prediction and then computes the negative gradients of the loss function for the current predictions. These negative gradients indicate the direction and magnitude of the steepest loss reduction.

At each iteration, a new decision tree is trained to predict these negative gradients (also called pseudo-residuals). For example, if the true CLV is \$2,000 but the model currently predicts \$1,500, the residual is +\$500. The next tree learns to predict this residual using the same input features, nudging the ensemble's prediction closer to the true value. After each step, the residuals are recalculated based on the updated predictions, focusing training on the remaining errors.

The **learning rate** scales each tree's contribution, preventing overfitting by ensuring gradual, cautious updates. Instead of correcting all errors at once, the model takes many small, informed steps toward a better solution.

The final prediction sums the initial estimate and all scaled tree outputs, each making targeted corrections. Through this stepwise refinement, gradient boosting transforms many simple models into a highly accurate predictor. Here is what this process looks like:

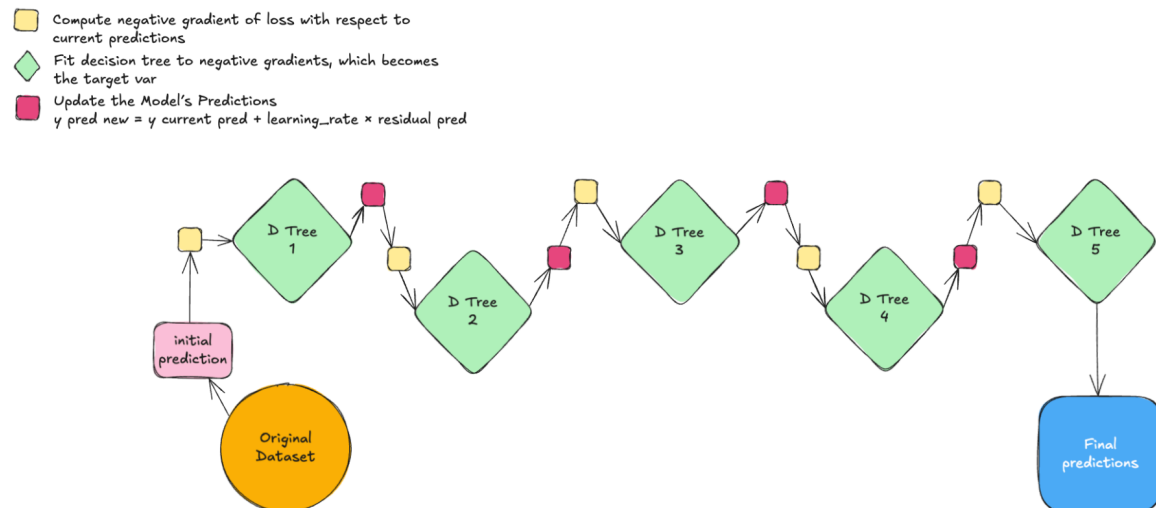


Figure 1

Can We Build It? Yes, We Can!

With Python, initializing and training a gradient boosting model is remarkably straightforward – it takes just two lines of code:

```
model = GradientBoostingClassifier(parameter_1,...,parameter_n)
```

or

```
model = GradientBoostingRegressor(parameter_1,...,parameter_n)
model.fit(X, y)
```

The first line sets up the model, and the second line uses the fit() function to learn the relationship between all the independent variables and CLV. While you can create a gradient boosting model using default settings, such as GradientBoostingRegressor(), sklearn offers several parameters to help control the tree's complexity and prevent overfitting. Here are a few common gradient boosting model-specific parameters:

1. `n_estimators`: Sets the number of trees in the forest.
2. `learning_rate`: Controls the speed of adjustments in the gradient descent.
3. `max_depth`: Limits the maximum height of the tree.

These parameters afforded to us by Python's Scikit-learn library make it pretty easy to control the gradient boosting model's input, size, and structure, but what's going on inside the 'black box'? In machine learning, when we say something is a 'black box', we mean that the internal process is hidden from us. Though sklearn makes it easy to use this function, we still want to understand what's going on inside.

Math in the Black Box

Our goal is to accurately predict the Customer Lifetime Value (CLV) of an insurance company's clients using a variety of predictors, including coverage type, income, and the number of policies. Since CLV is a continuous variable, we employ a regression gradient boosting model for this task.

How Regression Decision Trees Work

A regression decision tree estimates continuous outcomes by repeatedly splitting the dataset into smaller, more homogeneous groups based on the target values.

1. **Splitting**: At each node, the algorithm selects the feature and split point that leads to the greatest reduction in the sum of squared errors (SSE) within the resulting subsets.
2. **Leaf Assignment**: The Process continues recursively until a pre-set stopping criterion is met.
3. **Prediction**: The predicted value for any leaf is the average of the target values of the training samples in that group.

Let me illustrate this process:




Figure 2

Gradient Boosting Process

Gradient boosting works by building an ensemble of decision trees one at a time, where each new tree helps to correct the errors made by the combined predictions of the trees before it. The process begins with a baseline prediction for every data point: sklearn's function generally uses the mean of the training dataset's target variable. This provides a simple baseline estimate, serving as a starting point for the boosting process. This initial, feature-agnostic prediction establishes a starting point for the model.

Next, the algorithm evaluates the gradient of the loss function with respect to the current model's predictions. The loss function of our GradientBoostingRegressor is the Mean Squared Error loss function:

$$L_{MSE} = (y_i - \hat{y}_i)^2$$

To find the negative gradient, we must calculate the negative partial derivative of our loss function for the prediction (\hat{y}):

$$\begin{aligned}
-\frac{\partial L_{MSE}}{\partial \widehat{y}_i} &= -\frac{\partial (y_i - \widehat{y}_i)^2}{\partial \widehat{y}_i} \\
&= -2(y_i - \widehat{y}_i)
\end{aligned}$$

To focus on the gradient, let's omit the constant factor -2, which scales (adjusts the step size of) the gradient but doesn't affect the direction at all. We are left with:

$$(y_i - \widehat{y}_i)$$

Which is the expression for the residual. This tells us that in sklearn's GradientBoostingRegressor function, residuals and the negative gradient are actually the same. However, keep in mind that with other loss functions, the negative gradient may not be identical to the raw residuals as seen above.

Each tree is trained to predict this negative gradient. Before each new tree is built, the algorithm recalculates how much the current model's predictions need to change to most efficiently reduce the overall error, according to the chosen loss function. Once trained, the predictions of this tree are scaled by the learning rate parameter and added to the model's current predictions. This process is repeated iteratively: compute gradients → fit a new tree to them → scale its output → update the predictions, until a stopping criterion is met.

While studying this model, I would often refer to my notes on logistic regression to compare gradient descent and gradient ascent. I realized that gradient descent minimizes a loss function by moving in the direction of the negative gradient. In contrast, gradient ascent maximizes a function (like the log-likelihood in logistic regression) by moving in the direction of the positive gradient. Mathematically, one subtracts the gradient, the other adds it. With gradient descent, the goal is to find the minima (lowest point) of a loss function to reduce errors; with gradient ascent, the goal is to find the maxima (highest point) of a log-likelihood function to identify the most probable coefficients.

Computational Mathematics

With this context in mind, let's manually build a gradient boosting regressor. The dataset has already been cleaned and feature-engineered to feed the model. Be sure to check these sections out in the code file.

To manually build a gradient boosting regressor step by step, we start by defining helper functions that the tree will need. First up is a function that computes the sum of squared errors, which helps to evaluate how well a split explains the variance in the data. This function is the sum of all residuals squared. It can be written as:

```
'''  
  
def sse(y):  
    return ((y - y.mean()) ** 2).sum()  
'''
```

Next, we need to compute the negative gradient of the loss function, which for MSE is simply the residuals between the true and predicted values:

```
'''  
  
def compute_neg_gradient(y_true, y_pred):  
    return y_true - y_pred  
'''
```

Now, let's build a function, `best_sse_split`, that executes a regression decision tree's feature selection logic. This function sorts features and corresponding target residuals, examines thresholds halfway between consecutive unique feature values, evaluates the SSE for left and right splits at each threshold, and picks the threshold that yields the lowest combined SSE:

```
'''  
  
def best_sse_split(X_col, y):  
    sort_idx = np.argsort(X_col)  
    ...  
    if weighted_sse < best_sse:  
        best_sse = weighted_sse  
        best_threshold = threshold  
    return best_threshold, best_sse  
'''
```

Using `best_sse_split`, we recursively build the decision tree targeting residuals with `fit_tree`. The `fit_tree` function checks if any stopping criterion is met:

- The maximum allowed depth of the tree is reached (`max_depth == 0`).
- There are not enough samples to split (`len(X) < min_samples_split`).
- All target values in the node are identical (`np.std(y) == 0`).

If any of these are true, the node becomes a leaf. Its prediction is the mean of its current target values:

```
...  
  
def fit_tree(X, y, max_depth, min_samples_split):  
    if (max_depth == 0) or (X.shape[0] < min_samples_split) or  
    (np.std(y) == 0):  
        ...  
        left_subtree = fit_tree(X[left_mask], y[left_mask], max_depth - 1,  
min_samples_split)  
        right_subtree = fit_tree(X[right_mask], y[right_mask], max_depth - 1,  
min_samples_split)  
        return {  
            'leaf': False,  
            'feature': best_feature,  
            'threshold': best_threshold,  
            'left': left_subtree,  
            'right': right_subtree  
        }  
    ...
```

By recursively calling this function for each left and right branch, we are assembling the entire tree as a nested dictionary structure. Here is an example output:

```
'leaf': True,  
'feature': 19,  
'threshold': 0.5,  
'left': {'leaf': True, 'value': -5053.183893302541},  
'right': {'leaf': True, 'value': -4941.1732656999775}
```


We've built a function that implements a regression decision tree, but now we need the tree to make predictions. The goal is to enable our gradient boosting model to generate predictions by leveraging the outputs of its individual decision trees. The following function, `predict_tree`, handles this task for a manually constructed regression decision tree represented as a nested dictionary. Given a set of input samples, it returns the predicted values for each instance.

```
...  
  
def predict_tree(tree, X):  
    preds = np.zeros(X.shape[0])  
    for i, x in enumerate(X):  
        node = tree  
        while not node['leaf']:  
            if x[node['feature']] <= node['threshold']:  
                node = node['left']  
            else:  
                node = node['right']  
        preds[i] = node['value']  
    return preds  
...
```

For each sample, the function traverses the tree from the root to a leaf node, determining the prediction along the way. At each node, it checks whether it has reached a leaf; if not, it examines the feature and threshold that dictate the data split at that point. The function then compares the current sample's feature value to the threshold to decide which branch to follow next. This "walking" down the tree continues until a leaf node is reached, where the prediction corresponds to the average target value stored there. When all the samples have been processed, the function returns the `preds` array, which holds the tree's prediction for each input.

Finally, the master function `build_gradient_boosting` orchestrates boosting by iteratively fitting trees to residuals and updating predictions:

```

...
def build_gradient_boosting(X, y, n_estimators=50, learning_rate=0.1,
max_depth=10, min_samples_split=2):
    ...
    for i in range(n_estimators):
        residual = y - y_pred
        tree = fit_tree(
            X, residual, max_depth=max_depth,
min_samples_split=min_samples_split
        )
        update = predict_tree(tree, X)
        y_pred += learning_rate * update
        trees.append(tree)
    # Return model state as a dictionary
    return {
        'init_value': init_value,
        'trees': trees,
        'learning_rate': learning_rate,
        'max_depth': max_depth,
        'min_samples_split': min_samples_split
    }
...

```

It starts with an initial prediction equal to the mean of the targets and iteratively fits trees to the residual errors. Then, `build_gradient_boosting` updates current predictions with scaled outputs of each new tree. It stores all trees from the final model, along with the initial prediction, learning rate, maximum tree depth, and minimum samples required to split a node. Finally, our gradient boosting regressor is constructed.

Predictions

Now that we've built our gradient boosting model, let's create a function to predict the Customer Lifetime Value (CLV) for new insurance customers!

The prediction function starts by initializing an array where every sample's prediction is set to the initial baseline—usually the mean target value from the training set. It then sequentially processes each tree in the trained ensemble. For each tree, it predicts the residual values for all samples, scales those predictions by the learning rate, and adds the result to the running total of predictions. After iterating through all trees, the function outputs the final combined predictions for every input sample, effectively aggregating the incremental corrections each tree contributes.

```
...  
def predict_gradient_boosting(X, trees, init_val, learning_rate):  
    y_pred = np.full(X.shape[0], init_val)  
    for tree in trees:  
        y_pred += learning_rate * predict_tree(tree, X)  
    return y_pred  
...
```

Once you run this code, the gradient boosting regressor will have made its predictions.

Evaluation

With predictions in hand, it's time to evaluate how well our Gradient Boosting Regressor performed against the actual CLV values. Since this is a regression problem, we rely on metrics designed for continuous targets: Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), and the R-squared score. Together, these metrics help us understand the average prediction error, the impact of large deviations, and how well the model explains variance in the data.

Let's take a look at the results of the Gradient Boosting Regressor:

1. Mean Absolute Error (MAE): 689.18
2. Mean Squared Error (MSE): 3,963,623.45

3. Root Mean Squared Error (RMSE): 1,990.89
4. R-squared Score: 0.9231

To put these numbers in perspective, consider the distribution of the target variable:

- Maximum CLV: 83,325.38
- Average CLV: 8,004.94
- Minimum CLV: 1,898.01

These results tell us several things:

1. The MAE of 689.18 indicates that, on average, our model's predictions deviate from the true values by approximately \$689. This is a relatively low average error, especially compared to the average target value of around \$8,000, highlighting strong predictive accuracy.
2. The MSE of 3,963,623.45, which squares each prediction error before averaging, underscores that most errors are small but still appropriately penalizes larger deviations more heavily, reflecting the model's performance in the presence of outliers.
3. The RMSE of 1,990.89, being the square root of the MSE, reveals that on average, the size of the prediction errors is just under \$2,000, confirming a tight fit, especially given the natural variance in real-world data.
4. The R-squared score of 0.9231 demonstrates that our model explains over 92% of the variance in the target variable. This strong fit signifies that the model captures the key underlying patterns influencing predictions.

It is worth noting that, although the sklearn Gradient Boosting Regressor is often considered a strong benchmark, it performs slightly worse than our manual implementation on this dataset according to these metrics. This evaluation focuses on the manually built Gradient Boosting Regressor. Overall, these results demonstrate that our custom model not only matches but can even surpass standard implementations.

Interpretation

Based on our evaluation metrics, it's evident that the manually implemented Gradient Boosting Regressor offers a strong and reliable baseline for predicting Customer Lifetime Value (CLV). With a Mean Absolute Error (MAE) of just \$689, the model provides highly accurate predictions relative to the average customer value, demonstrating impressive precision. The Root Mean Squared Error (RMSE) of approximately \$1,991 further confirms that the majority of prediction errors are reasonably small, although some outliers still cause larger deviations, as reflected in the Mean Squared Error (MSE).

The R^2 score of 0.9231 indicates that the model explains over 92% of the variance in CLV, highlighting its effectiveness in capturing the key factors driving customer value. This represents a significant improvement over typical benchmarks and suggests that our model can robustly handle the complexities present in real-world customer behavior.

While the sklearn Gradient Boosting Regressor is a well-established model, it performs slightly worse on this dataset compared to our manual implementation, reinforcing the strength of this custom approach. However, as with any model, continuous refinement through feature engineering, hyperparameter tuning, or incorporating more robust loss functions could further enhance performance, especially in handling outliers for sensitive applications such as personalized pricing or risk assessments.

The feature importance analysis sheds light on the main drivers behind CLV predictions in our model. Key factors include higher education levels—particularly doctorate holders—correlating with increased customer value; variations in policy types and specialized insurance products impacting value distributions; retired employment status reflecting distinct spending habits; suburban residence revealing geographic influences; and vehicle class indicating different customer segments. These insights collectively empower more focused and data-driven business strategies, from targeted marketing campaigns to risk management and customer segmentation:

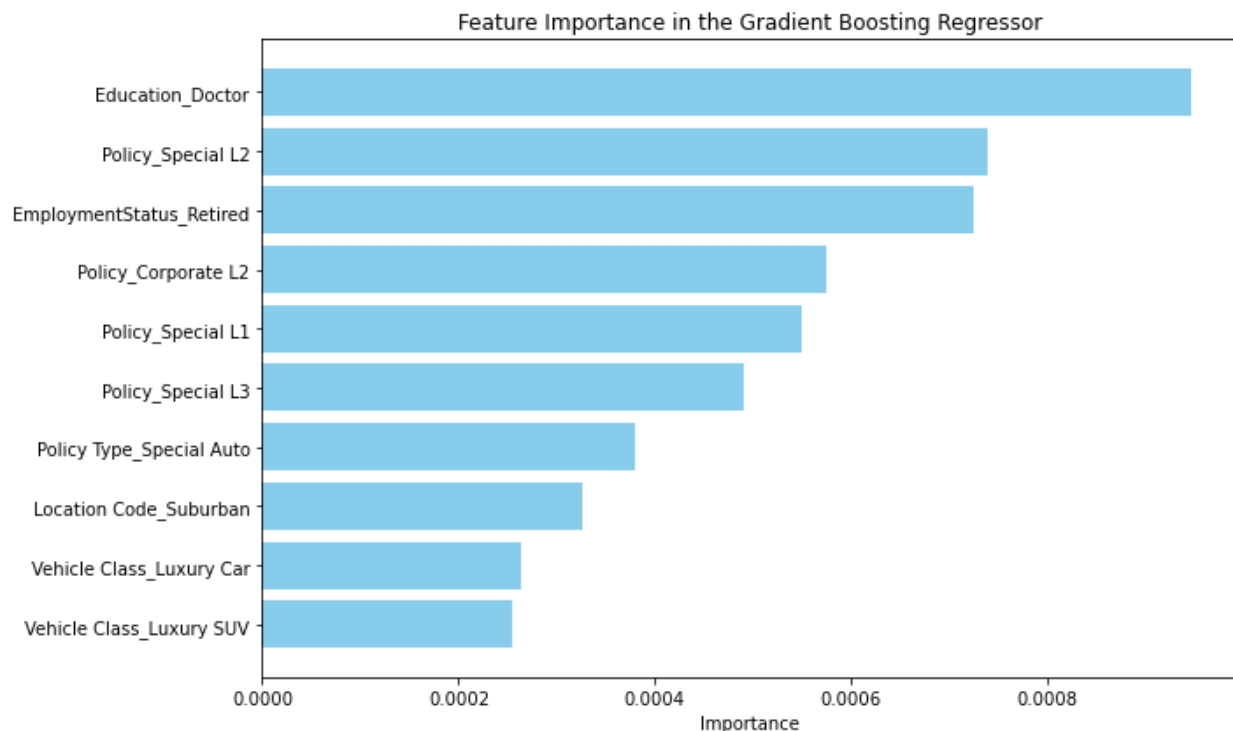


Figure 3

Conclusion

Gradient Boosting refines regression predictions by building a sequence of models that each learn from the mistakes of their predecessors. Unlike Random Forests, which average independently grown trees, Gradient Boosting trains new trees to correct residual errors sequentially, guided by gradient descent to minimize loss.

In this post, we moved from intuition to implementation, exploring how gradient boosting uses decision trees to fit residuals, leverages gradients for iterative updates, and incrementally reduces prediction error. We demonstrated training a Gradient Boosting Regressor on real CLV data, evaluated its performance, and interpreted its outputs.

Our evaluation shows the model explains 92% of CLV variance and maintains a manageable prediction error, handling most cases well despite some outliers. Gradient Boosting proves to be a powerful, flexible tool for practical business forecasting.

Ultimately, gradient boosting is about learning from errors; transforming residuals into improved predictions, one tree at a time. By understanding its core mechanics, you gain the insight needed to fine-tune, extend, and deploy it effectively.

In this post, we've learned:

- What gradient boosting machines are
- The structure of a gradient boosting machine
- How to build a gradient boosting regressor using the scikit-learn library in Python
- How to manually build a gradient boosting regressor in Python
- How to predict values using scikit and manually in Python
- How to evaluate the model using MAE, MSE, RMSE, and R-squared score
- How to interpret model evaluations
- As a bonus: Methods to optimize tree number and visualize feature influence

That's a wrap on gradient boosting machines! Next, we'll dive into XGBoost, an optimized, scalable variation of gradient boosting that takes the technique to the next level. Stay tuned!